

Distributed Algorithms & Protocols Cheat Sheet

Algorithm / Protocol	What It Does	Best Case	Worst Case
Rumor Mongering	Probabilistic gossip protocol spreading updates.	$O(\log n)$ rounds	Some nodes may not get update
Anti-Entropy (Push/Pull)	Nodes exchange data to reach consistency.	$O(\log n)$ time	$O(n)$ redundant messages
Napster	Centralized P2P with central index.	Fast lookup	Single point of failure
Gnutella	Fully decentralized P2P using flooding.	Simple, resilient	High traffic, not scalable
Pastry DHT	Prefix-based structured routing overlay.	$O(\log n)$ hops	$O(n)$ on failure
Freenet	Anonymous distributed datastore.	Fast for local keys	Slow for unpopular data
BitTorrent	File sharing with tit-for-tat piece exchange.	Parallel fast download	Slow if few seeders
Chord DHT	Consistent hashing for lookup.	$O(\log n)$ hops	$O(n)$ under churn
Lamport Clock	Logical timestamps for event ordering.	Constant	$O(n)$ sync overhead
Vector Clock	Tracks causality among events precisely.	$O(n)$ vector size	High memory usage
Cristian's Algorithm	Synchronizes client clock with server.	Low drift	Large drift on network delay
Chang-Roberts Election	Ring-based leader election using IDs.	$O(n)$	$O(n^2)$ messages
Ricart–Agrawala	Request-reply mutual exclusion.	$O(n)$	$O(n)$ messages per entry
Maekawa's Algorithm	Quorum-based mutual exclusion.	$O(\sqrt{n})$ messages	Deadlock risk
Suzuki–Kasami	Token-based mutual exclusion.	$O(1)$ entry delay	$O(n)$ to find token
Raymond's Tree Algorithm	Tree-based token passing mutual exclusion.	$O(\log n)$	$O(n)$ delay
GHS MST	Builds distributed minimum spanning tree.	$O(n \log n)$	$O(n^2)$ messages
FLP Impossibility	Consensus impossible in async system w/1 crash.	-	-
Raft Consensus	Leader-based replicated log agreement.	$O(1)$ commit	Delayed elections
Paxos	Fault-tolerant consensus protocol.	$O(1)$ after quorum	$O(n)$ retries
Byzantine Agreement	Achieves agreement despite malicious faults.	$3f+1$ nodes	$O(n^2)$ messages
2-Phase Commit	Atomic commit using coordinator.	Fast commit	Blocking on crash
3-Phase Commit	Adds pre-commit to avoid blocking.	Reliable → commit	Abort on failures
AFS	Distributed file system w/ whole-file caching.	Local access	Invalidation overhead
Coda FS	AFS extension w/ replication + offline ops.	Fault tolerant	Manual conflict fix
Amazon Dynamo	Eventually consistent key-value store.	High availability	Temporary inconsistency
Bigtable	Google's distributed storage system.	$O(1)$ lookup	$O(\log n)$ w/ compaction
Percolator	Incremental ACID processing over Bigtable.	Fast incremental update	High write contention
Corona	Pastry-based cluster monitoring system.	Balanced load	Slower under churn
PageRank	Ranks pages via incoming links.	Fast on sparse graphs	$O(n^2)$ dense graphs